

dlnd_face_generation

May 29, 2018

1 Face Generation

In this project, you'll use generative adversarial networks to generate new images of faces. ###
Get the Data You'll be using two datasets in this project: - MNIST - CelebA

Since the celebA dataset is complex and you're doing GANs in a project for the first time, we want you to test your neural network on MNIST before CelebA. Running the GANs on MNIST will allow you to see how well your model trains sooner.

If you're using [FloydHub](#), set `data_dir` to `"/input"` and use the [FloydHub data ID](#) `"R5KrjnANiKVhLWApXhNBe"`.

```
In [1]: data_dir = './data'
```

```
# FloydHub - Use with data ID "R5KrjnANiKVhLWApXhNBe"  
#data_dir = '/input'
```

```
"""  
DON'T MODIFY ANYTHING IN THIS CELL  
"""
```

```
import helper
```

```
helper.download_extract('mnist', data_dir)  
helper.download_extract('celeba', data_dir)
```

```
Downloading mnist: 9.92MB [00:01, 5.99MB/s]
```

```
Extracting mnist: 100%|| 60.0K/60.0K [00:08<00:00, 6.90KFile/s]
```

```
Downloading celeba: 1.44GB [02:20, 10.3MB/s]
```

```
Extracting celeba...
```

1.1 Explore the Data

1.1.1 MNIST

As you're aware, the [MNIST](#) dataset contains images of handwritten digits. You can view the first number of examples by changing `show_n_images`.

```
In [2]: show_n_images = 25
```

```
"""
DON'T MODIFY ANYTHING IN THIS CELL
"""


%matplotlib inline
import os
from glob import glob
from matplotlib import pyplot

mnist_jpegs = glob(os.path.join(data_dir, 'mnist/*.jpg'))
mnist_images = helper.get_batch(mnist_jpegs[:show_n_images], 28, 28, 'L')

pyplot.imshow(helper.images_square_grid(mnist_images, 'L'))
# helper.images_square_grid had to be adapted, i.e. convert image to RGBA,
# cmap = 'gray' unnecessary now.
# see also: https://github.com/matplotlib/matplotlib/issues/10616/
```

```
/usr/local/lib/python3.5/dist-packages/matplotlib/font_manager.py:279: UserWarning: Matplotlib is building the font cache using fc-list.
'Matplotlib is building the font cache using fc-list. '
```

```
Out[2]: <matplotlib.image.AxesImage at 0x7f44eec739e8>
```




1.1.2 CelebA

The [CelebFaces Attributes Dataset \(CelebA\)](#) dataset contains over 200,000 celebrity images with annotations. Since you're going to be generating faces, you won't need the annotations. You can view the first number of examples by changing `show_n_images`.

```
In [3]: show_n_images = 25
```

```
"""  
DON'T MODIFY ANYTHING IN THIS CELL  
"""  
celeba_jpegs = glob(os.path.join(data_dir, 'img_align_celeba/*.jpg'))  
mnist_images = helper.get_batch(celeba_jpegs[:show_n_images], 28, 28, 'RGB')  
pyplot.imshow(helper.images_square_grid(mnist_images, 'RGB'))
```

```
Out [3]: <matplotlib.image.AxesImage at 0x7f44e6aea940>
```



1.2 Preprocess the Data

Since the project's main focus is on building the GANs, we'll preprocess the data for you. The values of the MNIST and CelebA dataset will be in the range of -0.5 to 0.5 of 28x28 dimensional images. The CelebA images will be cropped to remove parts of the image that don't include a face, then resized down to 28x28.

The MNIST images are black and white images with a single [color channel]([https://en.wikipedia.org/wiki/Channel_\(digital_image%29\)](https://en.wikipedia.org/wiki/Channel_(digital_image%29))) while the CelebA images have [3 color channels (RGB color channel)]([https://en.wikipedia.org/wiki/Channel_\(digital_image%29#RGB_Images\)](https://en.wikipedia.org/wiki/Channel_(digital_image%29#RGB_Images))). ## Build

the Neural Network You'll build the components necessary to build a GANs by implementing the following functions below: - model_inputs - discriminator - generator - model_loss - model_opt - train

1.2.1 Check the Version of TensorFlow and Access to GPU

This will check to make sure you have the correct version of TensorFlow and access to a GPU

```
In [4]: """
        DON'T MODIFY ANYTHING IN THIS CELL
        """

        from distutils.version import LooseVersion
        import warnings
        import tensorflow as tf

        # Check TensorFlow Version
        assert LooseVersion(tf.__version__) >= LooseVersion('1.0'), \
            'Please use TensorFlow version 1.0 or newer. You are using {}'.format(tf.__version__)
        print('TensorFlow Version: {}'.format(tf.__version__))

        # Check for a GPU
        if not tf.test.gpu_device_name():
            warnings.warn('No GPU found. Please use a GPU to train your neural network.')
        else:
            print('Default GPU Device: {}'.format(tf.test.gpu_device_name()))
```

TensorFlow Version: 1.4.1

Default GPU Device: /device:GPU:0

1.2.2 Input

Implement the model_inputs function to create TF Placeholders for the Neural Network. It should create the following placeholders: - Real input images placeholder with rank 4 using image_width, image_height, and image_channels. - Z input placeholder with rank 2 using z_dim. - Learning rate placeholder with rank 0.

Return the placeholders in the following the tuple (tensor of real input images, tensor of z data)

```
In [5]: import problem_unittests as tests

        def model_inputs(image_width, image_height, image_channels, z_dim):
            """
            Create the model inputs
            :param image_width: The input image width
            :param image_height: The input image height
            :param image_channels: The number of image channels
            :param z_dim: The dimension of Z
            :return: Tuple of (tensor of real input images, tensor of z data, learning rate)
```

```

"""
# TODO: Implement Function
imreal = tf.placeholder(tf.float32, shape = (None, image_width,
                                             image_height, image_channels))

z      = tf.placeholder(tf.float32, shape = (None, z_dim))
lr      = tf.placeholder(tf.float32)
return imreal, z, lr

"""
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""
tests.test_model_inputs(model_inputs)

```

Tests Passed

1.2.3 Discriminator

Implement discriminator to create a discriminator neural network that discriminates on images. This function should be able to reuse the variables in the neural network. Use `tf.variable_scope` with a scope name of "discriminator" to allow the variables to be reused. The function should return a tuple of (tensor output of the discriminator, tensor logits of the discriminator).

```

In [6]: def discriminator(images, reuse=False):
        """
        Create the discriminator network
        : param images : Tensor of input image(s)
        : param reuse   : Boolean if the weights should be reused
        : return        : Tuple of (tensor output of the discriminator,
                                   tensor logits of the discriminator)
        """
        def convolution(in_, f, k = (5, 5), s = (2, 2), p = 'SAME',
                        activation = None, use_bias = False):
            return tf.layers.conv2d(in_, filters = f,
                                     kernel_size = k, strides = s,
                                     padding = p, data_format = 'channels_last',
                                     activation = activation, use_bias = use_bias)

        def fullyconnected(in_, num_units = 1, use_bias = False):
            return tf.layers.dense(in_, units = num_units,
                                    activation = None, use_bias = use_bias)

        def batchnorm(in_):
            return tf.layers.batch_normalization(in_, training = True)
        #training = True, as discriminator not used for inference

        def activation(in_, alpha=0.02, kind='leakyReLU'):
            if kind == 'leakyReLU':

```

```

        return tf.maximum(alpha*in_, in_)
    elif kind == 'sigmoid':
        return tf.sigmoid(in_)

# TODO: Implement Function
with tf.variable_scope("discriminator", reuse = reuse):
    layers = 3
    layer = images
    list_layers = list(range(layers))
    for i in list_layers:
        if i != list_layers[-1]:
            layer = convolution(layer, f=2*(i+1)*int(images.shape[1]))
            if i != 0:
                layer = batchnorm(layer)
            layer = activation(layer)
        else:
            reshape_dims = layer.shape[1]*layer.shape[2]*layer.shape[3]
            layer = tf.reshape(layer, shape = (-1, reshape_dims))
            logits = fullyconnected(layer)
            out = activation(logits, kind = 'sigmoid')

    return logits, out

"""
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""
tests.test_discriminator(discriminator, tf)

```

Tests Passed

1.2.4 Generator

Implement generator to generate an image using z . This function should be able to reuse the variables in the neural network. Use `tf.variable_scope` with a scope name of "generator" to allow the variables to be reused. The function should return the generated $28 \times 28 \times \text{out_channel_dim}$ images.

```

In [22]: def generator(z, out_channel_dim, is_train=True):
    """
    Create the generator network
    :param z: Input z
    :param out_channel_dim: The number of channels in the output image
    :param is_train: Boolean if generator is being used for training
    :return: The tensor output of the generator
    """
    def fullyconnected(in_, num_units):
        return tf.layers.dense(in_, units = num_units, activation = None,

```

```

        use_bias = False)

def transposed_convolution(in_, f, k = (5, 5), s = (2, 2), p = 'SAME',
                           activation = None, use_bias = False):
    return tf.layers.conv2d_transpose(in_, filters = f,
                                       kernel_size = k, strides = s,
                                       padding = p, data_format = 'channels_last',
                                       activation = activation, use_bias = use_bias)

def batchnorm(in_):
    return tf.layers.batch_normalization(in_, training = is_train)
#training = True, as discriminator not used for inference

def leakyReLU(in_, alpha=0.02):
    return tf.maximum(alpha*in_, in_)

def tanh(in_):
    return tf.tanh(in_)

# TODO: Implement Function
reuse = False if is_train else True
with tf.variable_scope("generator", reuse = reuse):
    ''' --- APPARENTLY TOO COMPLICATED ---
    in_dim = (-1, 3, 3, 256)
    layer = fullyconnected(z, num_units = in_dim[1]*in_dim[2]*in_dim[3])

    # VALID padding and stride = 1 to generate more layers
    # (3,3,256) -> (5, 5, 128) - kernel size k = 3

    layer = tf.reshape(layer, shape = in_dim)
    layer = transposed_convolution(layer, f = 128, s = (1, 1), k = (3, 3), p = 'VALID')

    layer = batchnorm(layer)
    layer = leakyReLU(layer)

    # (5, 5, 64) -> (7, 7, 32) - k = 3
    layer = transposed_convolution(layer, f = 64, s = (1, 1), k = (3, 3), p = 'VALID')

    layer = batchnorm(layer)
    layer = leakyReLU(layer)

    # (7,7,64) -> (10, 10, 32) - k = 4
    layer = transposed_convolution(layer, f = 32, s = (1, 1), k = (4, 4), p = 'VALID')

    layer = batchnorm(layer)
    layer = leakyReLU(layer)

```

```

# (10, 10, 32) -> (14, 14, 16) - k = 5
layer = transposed_convolution(layer, f = 16, s = (1, 1), p = 'VALID')
#layer = transposed_convolution(layer, f = 16)

layer = batchnorm(layer)
layer = leakyReLU(layer)

# (14, 14, 16) -> (21, 21, 8) - k = 8
layer = transposed_convolution(layer, f = 8, s = (1, 1), k = (8, 8), p = 'VALID')

layer = batchnorm(layer)
layer = leakyReLU(layer)

# (21, 21, 9) -> (28,28,3) - k = 8
layer = transposed_convolution(layer, f = out_channel_dim, s = (1, 1), k = (8,
#layer = transposed_convolution(layer, f = out_channel_dim)

# layer = batchnorm(layer)
layer = tanh(layer)/2
# leakyReLU(layer); to get values in [-0.5, 0.5] as tanh gives [-1, 1]
'''

# --- INPUT LAYER ---
# Dimension of (3,3,64) reshaped into matrix form (batch_size, 3*3*64)
in_dim = (-1, 3, 3, 64)
layer = fullyconnected(z, num_units = in_dim[1]*in_dim[2]*in_dim[3])
layer = tf.reshape(layer, shape = in_dim)

# (3, 3, 64) -> (7, 7, 32)
layer = transposed_convolution(layer, f = 32, s = (1, 1), k = (5, 5), p = 'VALID')
layer = batchnorm(layer)
layer = leakyReLU(layer, alpha = 0.2)

# (7,7,32) -> (14, 14, 16)
layer = transposed_convolution(layer, f = 16)
layer = batchnorm(layer)
layer = leakyReLU(layer, alpha = 0.2)

# (14, 14, 16) -> (28,28,1 or 3)
layer = transposed_convolution(layer, f = out_channel_dim)
layer = tanh(layer) / 2
return layer

'''
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
'''

tests.test_generator(generator, tf)

```


Tests Passed

1.2.5 Loss

Implement `model_loss` to build the GANs for training and calculate the loss. The function should return a tuple of (discriminator loss, generator loss). Use the following functions you implemented: - `discriminator(images, reuse=False)` - `generator(z, out_channel_dim, is_train=True)`

```
In [39]: def model_loss(input_real, input_z, out_channel_dim):
        """
        Get the loss for the discriminator and generator
        :param input_real: Images from the real dataset
        :param input_z: Z input
        :param out_channel_dim: The number of channels in the output image
        :return: A tuple of (discriminator loss, generator loss)
        """
        smooth = 1
        # TODO: Implement Function
        generator_out = generator(input_z, out_channel_dim, is_train = True)
        d_logits_real, d_out_real = discriminator(input_real, reuse = False)
        d_logits_fake, d_out_fake = discriminator(generator_out, reuse = True)

        d_loss_real = tf.reduce_mean\
        (tf.nn.sigmoid_cross_entropy_with_logits\
        (labels = tf.ones_like(d_out_real)*smooth, logits = d_logits_real, \
        name = 'd_loss_real'))

        d_loss_fake = tf.reduce_mean\
        (tf.nn.sigmoid_cross_entropy_with_logits\
        (labels = tf.zeros_like(d_out_fake), logits = d_logits_fake, \
        name = 'd_loss_fake'))

        d_loss = d_loss_real + d_loss_fake
        g_loss = tf.reduce_mean\
        (tf.nn.sigmoid_cross_entropy_with_logits\
        (labels=tf.ones_like(d_out_fake), logits = d_logits_fake, name = 'generator_loss'))
        return (d_loss, g_loss)

        """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """
        tests.test_model_loss(model_loss)
```

Tests Passed

1.2.6 Optimization

Implement `model_opt` to create the optimization operations for the GANs. Use `tf.trainable_variables` to get all the trainable variables. Filter the variables with names that are in the discriminator and generator scope names. The function should return a tuple of (discriminator training operation, generator training operation).

```
In [40]: def model_opt(d_loss, g_loss, learning_rate, beta1):
        """
        Get optimization operations
        :param d_loss: Discriminator loss Tensor
        :param g_loss: Generator loss Tensor
        :param learning_rate: Learning Rate Placeholder
        :param beta1: The exponential decay rate for the 1st moment in the optimizer
        :return: A tuple of (discriminator training operation, generator training operation)
        """

        # TODO: Implement Function
        t_var = tf.trainable_variables ()
        d_var = [d for d in t_var if d.name.startswith("discriminator")]
        g_var = [g for g in t_var if g.name.startswith("generator")]

        update_ops = tf.get_collection(tf.GraphKeys.UPDATE_OPS)
        with tf.control_dependencies(update_ops):
            d_opt = tf.train.AdamOptimizer(learning_rate = learning_rate, beta1 = beta1)\
                .minimize(d_loss, var_list = d_var)
            g_opt = tf.train.AdamOptimizer(learning_rate = learning_rate, beta1 = beta1)\
                .minimize(g_loss, var_list = g_var)
        return (d_opt, g_opt)

        """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """

        tests.test_model_opt(model_opt, tf)
```

Tests Passed

1.3 Neural Network Training

1.3.1 Show Output

Use this function to show the current output of the generator during training. It will help you determine how well the GANs is training.

```
In [41]: """
        DON'T MODIFY ANYTHING IN THIS CELL
        """

        import numpy as np
```

```
def show_generator_output(sess, n_images, input_z, out_channel_dim, image_mode):
    """
    Show example output for the generator
    :param sess: TensorFlow session
    :param n_images: Number of Images to display
    :param input_z: Input Z Tensor
    :param out_channel_dim: The number of channels in the output image
    :param image_mode: The mode to use for images ("RGB" or "L")
    """
    cmap = None if image_mode == 'RGB' else 'gray'
    z_dim = input_z.get_shape().as_list()[-1]
    example_z = np.random.uniform(-1, 1, size=[n_images, z_dim])

    samples = sess.run(
        generator(input_z, out_channel_dim, False),
        feed_dict={input_z: example_z})

    images_grid = helper.images_square_grid(samples, image_mode)
    pyplot.imshow(images_grid, cmap=cmap)
    pyplot.show()
```

1.3.2 Train

Implement train to build and train the GANs. Use the following functions you implemented:

- model_inputs(image_width, image_height, image_channels, z_dim)
- model_loss(input_real, input_z, out_channel_dim)
- model_opt(d_loss, g_loss, learning_rate, beta1)

Use the show_generator_output to show generator output while you train. Running show_generator_output for every batch will drastically increase training time and increase the size of the notebook. It's recommended to print the generator output every 100 batches.

```
In [42]: import scipy.stats as st
def train(epoch_count, batch_size, z_dim,
          learning_rate, beta1, get_batches,
          data_shape, data_image_mode):
    """
    Train the GAN
    :param epoch_count: Number of epochs
    :param batch_size: Batch Size
    :param z_dim: Z dimension
    :param learning_rate: Learning Rate
    :param beta1: The exponential decay rate for the 1st moment in the optimizer
    :param get_batches: Function to get batches
    :param data_shape: Shape of the data
    :param data_image_mode: The image mode to use for images ("RGB" or "L")
    """
    # TODO: Build Model
    im_real, in_z, lr = model_inputs(data_shape[1], data_shape[2],
```

```

                                data_shape[3], z_dim)
d_loss, g_loss = model_loss(im_real, in_z, data_shape[3])
d_opt, g_opt    = model_opt(d_loss, g_loss, lr, beta1)

count = 0

with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())
    for epoch_i in range(epoch_count):
        for batch_images in get_batches(batch_size):
            # TODO: Train Model
            count += 1

            mu, sigma = 0, 0.25
            minval     = -0.5
            maxval     = 0.5
            rand_numbers = st.truncnorm.rvs((minval-mu)/sigma, (maxval-mu)/sigma,
                                           loc = mu, scale = sigma, size = (batch_size))
            train_d_loss, train_g_loss = sess.run([d_loss, g_loss, d_opt, g_opt], \
                                                  feed_dict = {im_real: batch_image,
                                                             in_z      : rand_numbers,
                                                             lr       : learning_rate})

            #Print losses, optimizer, generator
            if count % 200 == 0:
                print('Summary at {} batches'.format(count))
                print('discriminator model loss = {}'.format(train_d_loss))
                print('generator model loss = {}'.format(train_g_loss))

            show_generator_output(sess, 9, in_z, data_shape[3], data_image_mode='grayscale')

```

1.3.3 MNIST

Test your GANs architecture on MNIST. After 2 epochs, the GANs should be able to generate images that look like handwritten digits. Make sure the loss of the generator is lower than the loss of the discriminator or close to 0.

```

In [43]: %%time
         batch_size    = 128
         z_dim         = 100
         learning_rate  = 1e-4
         beta1         = 0.9

         """
         DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
         """

         epochs = 2


```

```
mnist_dataset = helper.Dataset('mnist', \
                                glob(os.path.join(data_dir, 'mnist/*.jpg')))
with tf.Graph().as_default():
    train(epochs, batch_size, z_dim, learning_rate, beta1, mnist_dataset.get_batches,
          mnist_dataset.shape, mnist_dataset.image_mode)
```

Summary at 200 batches

discriminator model loss = 0.11696434020996094


generator model loss = 2.6855640411376953




Summary at 400 batches

discriminator model loss = 0.04093387722969055

generator model loss = 3.930612564086914




Summary at 600 batches
discriminator model loss = 0.030817419290542603
generator model loss = 4.29692268371582



Summary at 800 batches

discriminator model loss = 0.036977220326662064

generator model loss = 4.170480251312256



CPU times: user 47 s, sys: 9.84 s, total: 56.8 s

Wall time: 54.7 s

1.3.4 CelebA

Run your GANs on CelebA. It will take around 20 minutes on the average GPU to run one epoch. You can run the whole epoch or stop when it starts to generate realistic faces.

```
In [44]: %%time
         batch_size    = 128
         z_dim         = 100
         learning_rate = 1e-4
         beta1         = 0.9

         """
         DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
         """
         epochs = 1
```

```

celeba_dataset = helper.Dataset('celeba', \
                                glob(os.path.join(data_dir, \
                                                    'img_align_celeba/*.jpg')))


with tf.Graph().as_default():
    train(epochs, batch_size, z_dim, learning_rate, beta1,
          celeba_dataset.get_batches, celeba_dataset.shape,
          celeba_dataset.image_mode)

```

Summary at 200 batches

discriminator model loss = 0.03729106858372688


generator model loss = 4.4327392578125



Summary at 400 batches

discriminator model loss = 0.0445941798388958


generator model loss = 4.771548748016357



Summary at 600 batches

discriminator model loss = 0.03814772889018059


generator model loss = 4.702086448669434



Summary at 800 batches

discriminator model loss = 0.06071583181619644

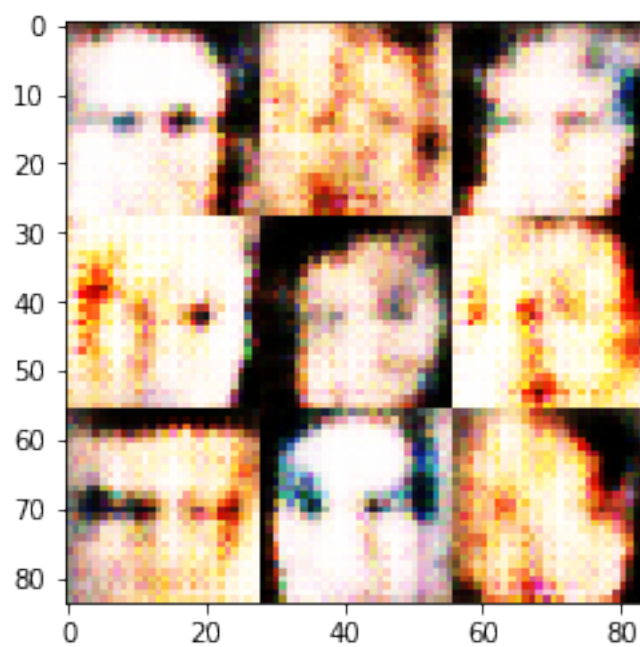
generator model loss = 4.194107532501221




Summary at 1000 batches

discriminator model loss = 0.047982364892959595

generator model loss = 4.353063583374023




Summary at 1200 batches
discriminator model loss = 0.055525705218315125
generator model loss = 4.489264488220215



Summary at 1400 batches

discriminator model loss = 0.061971329152584076

generator model loss = 4.829124450683594



CPU times: user 4min 6s, sys: 21.8 s, total: 4min 28s

Wall time: 4min 34s

1.3.5 Submitting This Project

When submitting this project, make sure to run all the cells before saving the notebook. Save the notebook file as "dlnd_face_generation.ipynb" and save it as a HTML file under "File" -> "Download as". Include the "helper.py" and "problem_unittests.py" files in your submission.